


## Preprocessing Pipeline

# Data Preprocessing: The Engine of AI

From Cleaning Tables to Training LLMs

Fuel



Imbalanced data  
analysis / treatment

Dimensionality  
reduction  
analysis / treatment

Data Transformation  
analysis / treatment

# Workshop Overview

Our journey through data preprocessing fundamentals to advanced LLM techniques

## 1 The Foundation

General data preprocessing for tabular and image data



## 2 The Pivot

Transition from traditional NLP to LLM preprocessing



## 3 Deep Dive - LLM Data Pipeline

Ingestion, filtering, deduplication, and PII redaction



## 4 Tokenization & Instruction Tuning

Subword tokenization and chat formatting techniques




## 5 Conclusion & Tools

Key takeaways and modern preprocessing frameworks



# Garbage In, Garbage Out

- ➊ 80% of an ML project is data preparation
- ➋ Model architecture **cannot fix** broken data
- ➌ "Data is **code**"



# Cleaning Structured Data

## Missing Values

Imputation (Mean/Median/KNN) vs. Dropping

## Outliers

Detection via Z-Score or IQR

## Scaling


Standardization vs. Normalization (Crucial for gradient descent)

Logistic Regression, SVM, PCA

distance-based algorithms (KNN)

## Practical Examples

### Before/After: Data Distribution Centering



# Preprocessing for Computer Vision

## Normalization

Rescaling pixel values **0-255 → 0-1**


## Resizing

**Fixed dimensions** for batch processing

## Augmentation


**Rotation, Flip, Zoom** (Fighting overfitting)

Image Augmentation Variations



# Preprocessing for Computer Vision

## OCR Pattern Matching Process




1 Identifying and Isolating words

2 Isolate letter contour from contour image

3 Process letters according to trained OCR input

4 Consolidate predictions according to OCR model

# Preprocessing for Computer Vision



# Traditional NLP vs. LLMs

## Traditional (TF-IDF/LSTM)

- Remove stopwords
- ✗ Stem words
- ↓ Lowercase everything
- ☰ Focus on keywords


## LLMs (Transformers)

- + Keep the context!
- ▲ Case matters
- = Stopwords matter
- ⌚ Focus on language structure



**Goal:** We want the model to learn language structure, not just keywords

# From Raw Web to Pre-training



## Key Data Sources



Common Crawl

GitHub

Wikipedia


Books

Forums

News Articles

# Quality Control at Scale

- 🌐 Language ID**  
Filtering out **wrong languages** (e.g., keeping only English/Persian)
- 📊 Metrics**  
**Perplexity scores**, text length, symbol-to-word ratio
- 刪 Removal**  
**HTML tags**, "lorem ipsum", toxic content



# Why Deduplication Matters?

## ⚠ The Problem

Duplicate data causes **memorization**, not **generalization**


## Exact Match

Removing identical strings using **SHA-256**

## Fuzzy Deduplication

**MinHash & LSH** for similar documents

## Data Leakage Prevention



# PII Redaction

## PII Must Be Scrubbed

**Personally Identifiable Information** removal is critical for privacy compliance

## Detecting

Emails, IP addresses, Phone numbers, Keys/Passwords

## Techniques

Regex patterns + Named Entity Recognition (NER)

## PII Redaction Example

Hello, my name is John Smith and I work at TechCorp. You can reach me at <EMAIL> or call my office at <PHONE>. Our company is located at 123 Main Street, New York, NY 10001. My employee ID is <ID> and my IP address is <IP>. Please don't share my password: <PASSWORD>.



PII Replaced with Tags

# How LLMs Read

## Not Just Splitting

It is **NOT** just splitting by space

## Sub-word Tokenization

BPE (Byte Pair Encoding) & **WordPiece**

## Efficiency

**Common words** = 1 token

**Rare words** = multiple tokens

## Sub-word Tokenization Example

**Unbelievable** ↓

**Un**  
Token ID: 1234

**##believe**  
Token ID: 5678

**##able**  
Token ID: 9012

# Instruction & Chat Formats

## Pre-training

Raw Text prediction

## Fine-tuning

Chat Template structure

## Format

Structured conversation with roles

## Masking

Loss calculated only on Assistant's response

## JSON to ChatML Conversion

### Original JSON

```
{ "instruction": "what is data preprocessing?",  
  "input": "", "output": "Data preprocessing is the  
  process of transforming raw data into an  
  understandable format." }
```



### ChatML Format

```
<|im_start|>user What is data preprocessing?  
<|im_end|> <|im_start|>assistant Data preprocessing
```

# Beyond Basic Tokenization

## SentencePiece

Unsupervised text tokenizer for language-agnostic tokenization

## Special Tokens

BOS, EOS, PAD, UNK - Control model behavior

## Vocabulary Size Impact

Larger vocab = fewer tokens per word but more parameters

## Tokenization Algorithm Comparison

Algorithm	Input Text	Output Tokens
WordPiece	Tokenization	Token ##ization
BPE	Tokenization	Token ization
SentencePiece	Tokenization	To ken ization
Unigram	Tokenization	Tok en iz ation

# Beyond Basic Tokenization

OOV = Out Of Vocabulary

## SentencePiece

**Unsupervised** text tokenizer for language-agnostic tokenization

## Special Tokens

**BOS, EOS, PAD, UNK** - Control model behavior

## Vocabulary Size Impact

**Larger vocab** = fewer tokens per word but more parameters

## Tokenization Algorithm Comparison

Feature	WordPiece	Byte-Pair Encoding (BPE)	Unigram
<b>Developer</b>	Google	Philip Gage (originally for compression)	Google (used by SentencePiece)
<b>Core Principle</b>	Segmentation based on <b>maximum likelihood</b> of words being split into components that already exist in the vocabulary (like morphemes).	Iteratively merging the <b>most frequent pair</b> of bytes/characters/tokens until the desired vocabulary size is reached.	Chooses the segmentation that <b>maximizes the probability</b> of the sequence based on a language model (often trained using the Expectation-Maximization algorithm).
<b>Typical Separator</b>	## prefix on subwords (e.g., run + ##ning ).	Underscore _ or a special hidden space byte (e.g., _token + iz + ation ).	Leading space _ or special token to denote the start of a word.
<b>Handling of OOV Words</b>	Excellent. Splits the unknown word into known root forms, prefixes, and suffixes.	Excellent. Guaranteed to break down any word into known subwords or individual characters/bytes.	Excellent. Offers multiple segmentation options and selects the most probable one.
<b>Common LLMs</b>	BERT, DistilBERT, ELECTRA	GPT-2, GPT-3, Llama 1/2, RoBERTa	Llama 3, XLNet, ALBERT, T5 (with SentencePiece)
<b>Farsi Example</b> (پژن)	پژن## + پژن (Segmentation based on known parts)	_پژن + پژن (Merging frequent pairs)	پ + پژن (Choosing the most probable sequence)

# Beyond Basic Fine-Tuning

## Supervised Fine-Tuning

SFT - Direct training on labeled examples

## RLHF

Reinforcement Learning from Human Feedback

## DPO

Direct Preference Optimization - Simpler than RLHF

## Parameter-Efficient

LoRA, QLoRA - Fine-tune with fewer resources

## Fine-Tuning Approaches

Method	One-Line Explanation	Primary Benefit	Need for Additional Data/Model	Complexity
SFT	Training on labeled examples to specialize in a task.	Simple, reliable, and well-understood.	Requires high-quality labeled (input/output) data.	Low
RLHF	Uses human feedback to train a Reward Model, then uses RL for alignment.	Aligns the model with complex human preferences.	Requires comparative data and a separate Reward Model.	High
DPO	Directly tunes the model based on human preferences without a Reward Model.	Simpler and more stable than RLHF.	Requires comparative data (preferred/rejected pairs).	Medium
PEFT (LoRA/QLoRA)	Fine-tuning by training a small number of parameters to drastically reduce resource needs.	Significant reduction in computational and memory resources.	Needs SFT data (but more efficient training).	Medium

# The Modern Tech Stack

## { } General Data Processing



Pandas

Tabular data manipulation



NumPy

Numerical computing



OpenCV

Computer vision

## ⚙️ LLM Specific



Hugging Face Datasets

Dataset management



Datatrove

Data processing pipeline



Text-Dedup

Text deduplication

# Scaling Data Processing



## Distributed Processing

Spark, Dask for handling massive datasets



## Data Sharding Strategies

Horizontal partitioning for parallel processing



## Quality Assessment

Automated metrics for data validation



## Automated Pipelines

End-to-end data cleaning workflows

## Distributed Data Processing Architecture



### Data Sources



### Distributed Storage



### Preprocessing



### Deduplication



### PII Removal



### Tokenization



### Formatting



### Processed Dataset

# Summary

↗ Data quality defines the **ceiling of performance**

💡 LLMs require "**Fuzzy**" cleaning, not destructive cleaning

☒ **Deduplication** is the most high-impact step for LLMs

## REAL WORLD

An example of Data Pre Processing  
(But not for models, In life)

### Data Preprocessing Checklist

- ✓ Clean tabular data (missing values, outliers)
- ✓ Normalize and augment image data
- ✓ Filter and clean web data for LLMs
- ✓ Remove duplicates (exact and fuzzy)

## **Review - Importance of Data Cleaning**

Data Cleaning is no longer a secondary step in data-driven projects; it is an indispensable part of their success.

Data that is not clean is unreliable, and making decisions based on it is useless or even harmful.

The reality is that the quality of the output of any analytical process is directly dependent on the quality of its input data.

Many data science professionals believe that without cleaning, data analysis will have no meaning or value.

Even the best algorithms cannot extract a correct result from incorrect data.

In the world of data, no analysis is reliable without clean data.

Clean Data = Correct Decision.

Data cleaning is a long-term investment for reducing direct and indirect costs.

# What Works and What Doesn't

## Best Practices

### Start Small

Begin with a **subset** of data before scaling

### Validate at Each Step

Check **data quality** after each transformation

### Document Your Pipeline

Create **reproducible** preprocessing workflows

### Iterate and Improve

Continuously **refine** based on model performance

## Common Pitfalls

### Over-cleaning Data

Removing **too much** context for LLMs

### Ignoring Data Drift

Not monitoring **changes** in data distribution

### Not Validating Steps

Skipping **quality checks** between stages

### Neglecting Deduplication

Underestimating its **impact** on model performance

# Questions & Discussion

## Key Resources

-  Hugging Face Documentation
-  Data Preprocessing Courses
-  GitHub Repository



## Ryan Heida

-  ryan@ryanheida.com / mr.ryanheida@gmail.com
-  www.ryanheida.com / links.ryanheida.com
-  t.me/RyanHeida  
[https://t.me/ML\\_Topics](https://t.me/ML_Topics)



## Thank You!

Scan for workshop feedback

